Digital Humans

Zhizheng Liu and Joe Lin

Research Goal

- Build **intelligent agents** that understand and interact with humans in outdoor environments (especially in urban contexts)
  - **Problem: Lack of human-scene and human-agent interaction data**

Scene

Human

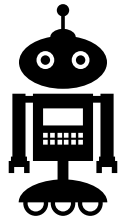

Agent

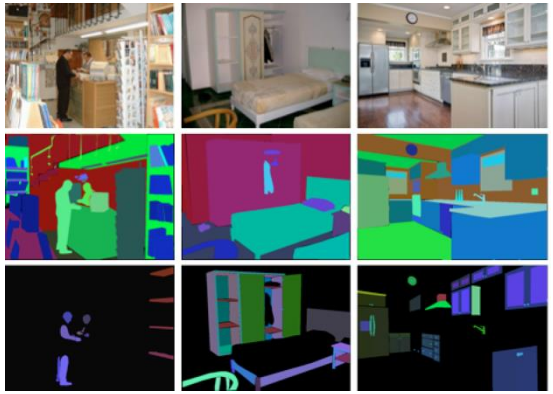

ADE20k

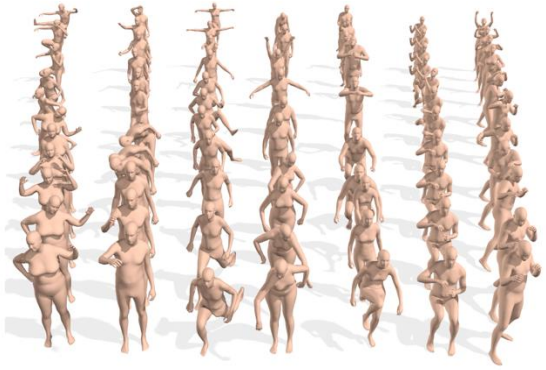

AMASS

MetaDrive + MetaUrban

Q: How do we get human-scene interaction data and learn from it?

A: CityWalkers and PedGen

Learning to Generate Diverse Pedestrian Movements from Web Videos with Noisy Labels

Zhizheng Liu, Joe Lin, Wayne Wu, Bolei Zhou

Representing Humans with SMPL

- Parameterized encoding of a 3D human mesh
  - Start with mean template shape $\bar{\mathbf{T}}$
  - Body **shape** parameters $\vec{\beta}$ (vector elements represent for ex. height, weight)
  - Body **pose** parameters $\vec{\theta}$ (rotation of each joint)
  - Global **orientation** and **translation** (pose in world space)

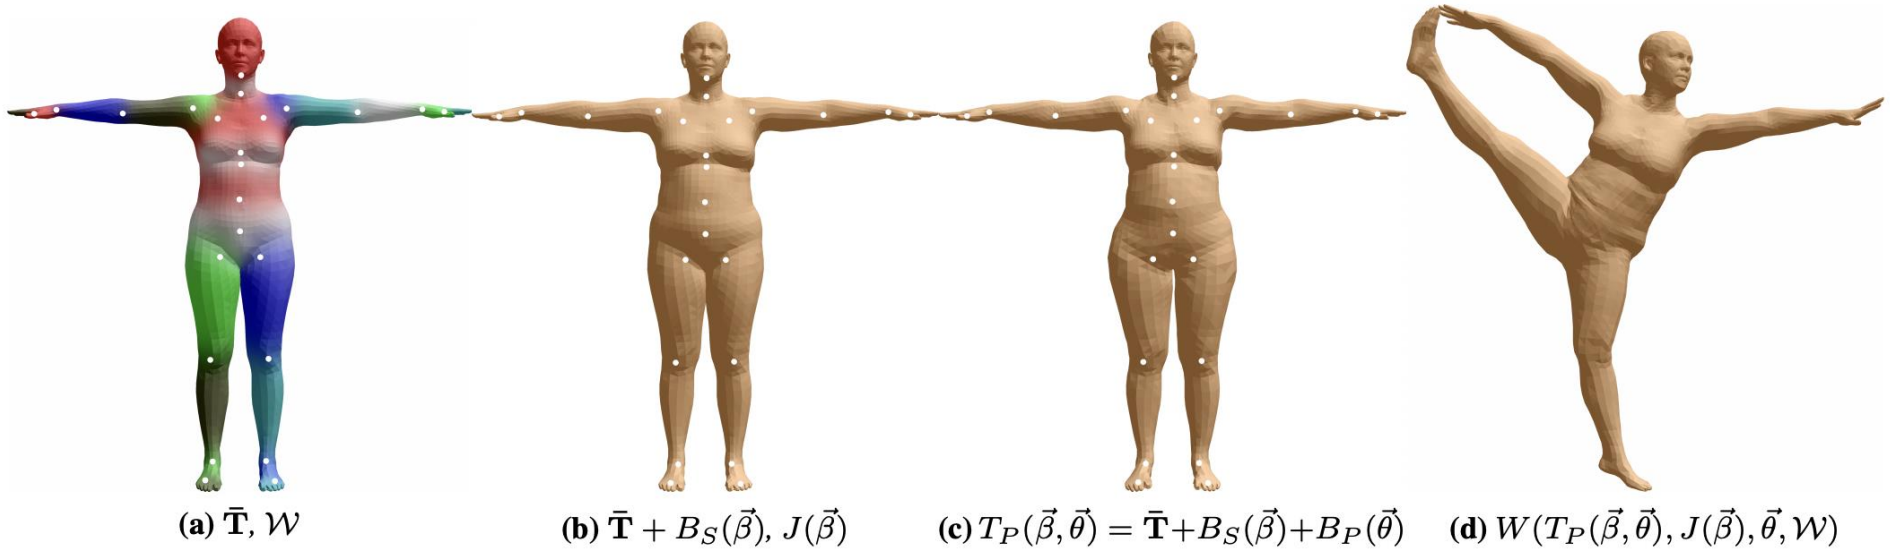

(a) $\bar{\mathbf{T}}, \mathcal{W}$ (b) $\bar{\mathbf{T}} + B_S(\vec{\beta}), J(\vec{\beta})$ (c) $T_P(\vec{\beta}, \vec{\theta}) = \bar{\mathbf{T}} + B_S(\vec{\beta}) + B_P(\vec{\theta})$ (d) $W(T_P(\vec{\beta}, \vec{\theta}), J(\vec{\beta}), \vec{\theta}, \mathcal{W})$

SMPL: A Skinned Multi-Person Linear Model. SIGGRAPH Asia 2015.

CityWalkers: Capturing Diverse Real-World Pedestrian Movements

- 30.8 hours of high-quality web videos with human motion pseudo-labels (SMPL)
- Diverse individual characteristics, human movements, urban environments, and route goals

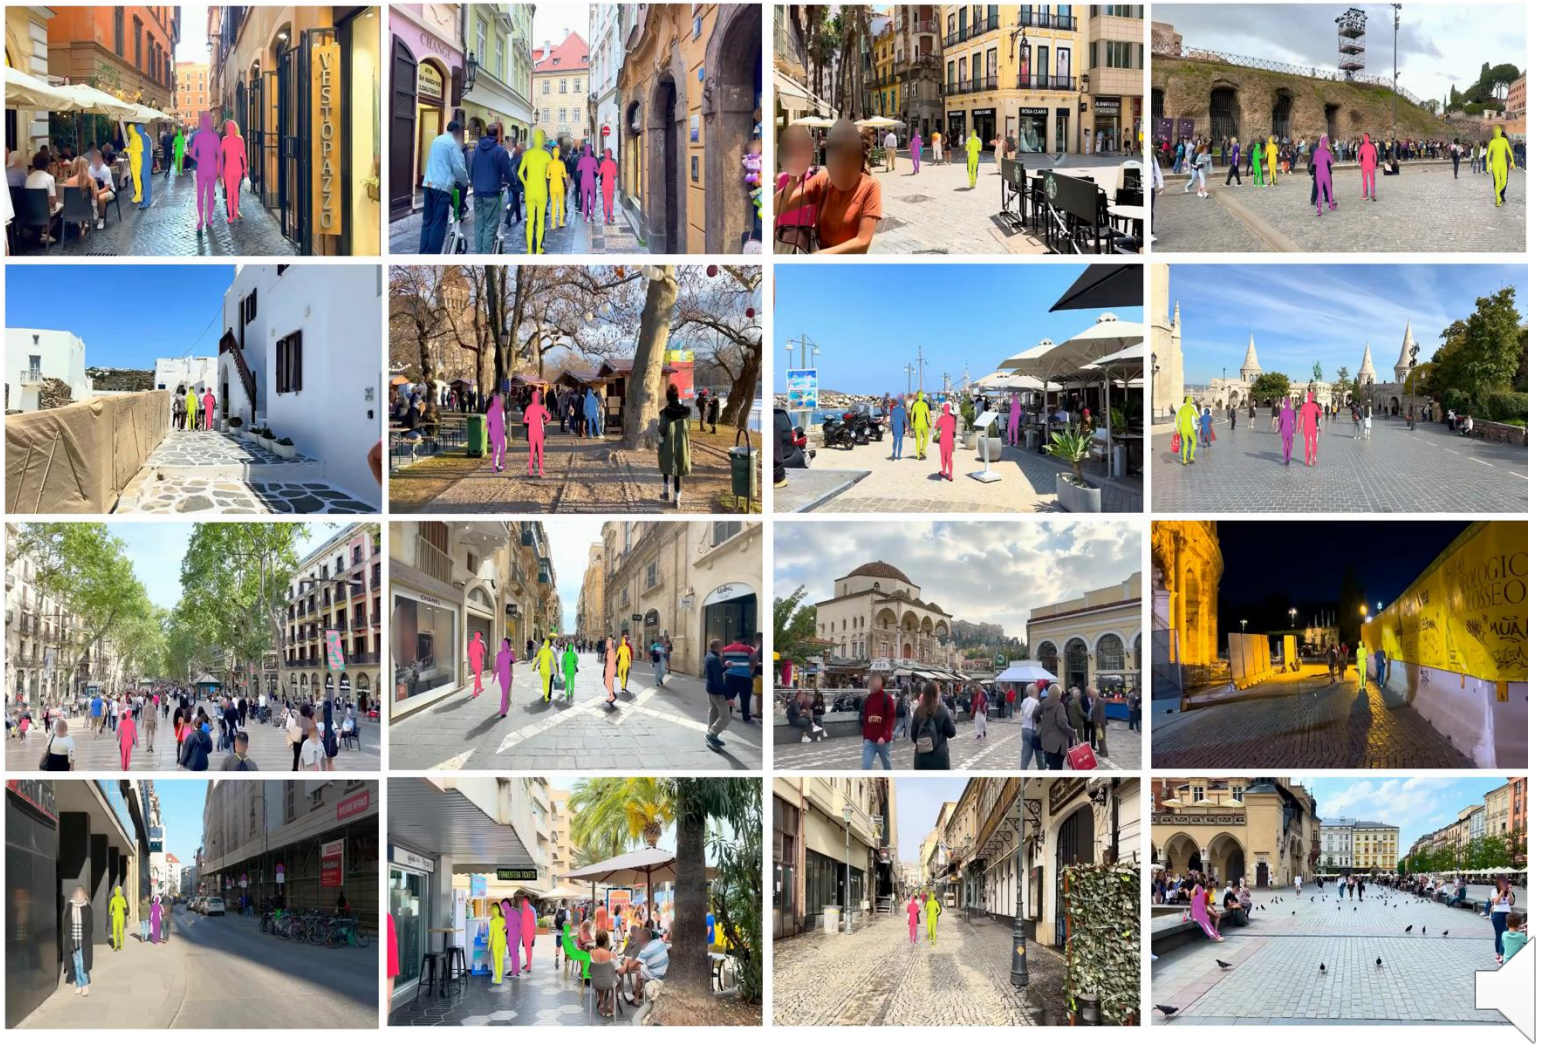

CityWalkers: Capturing Diverse Real-World Pedestrian Movements

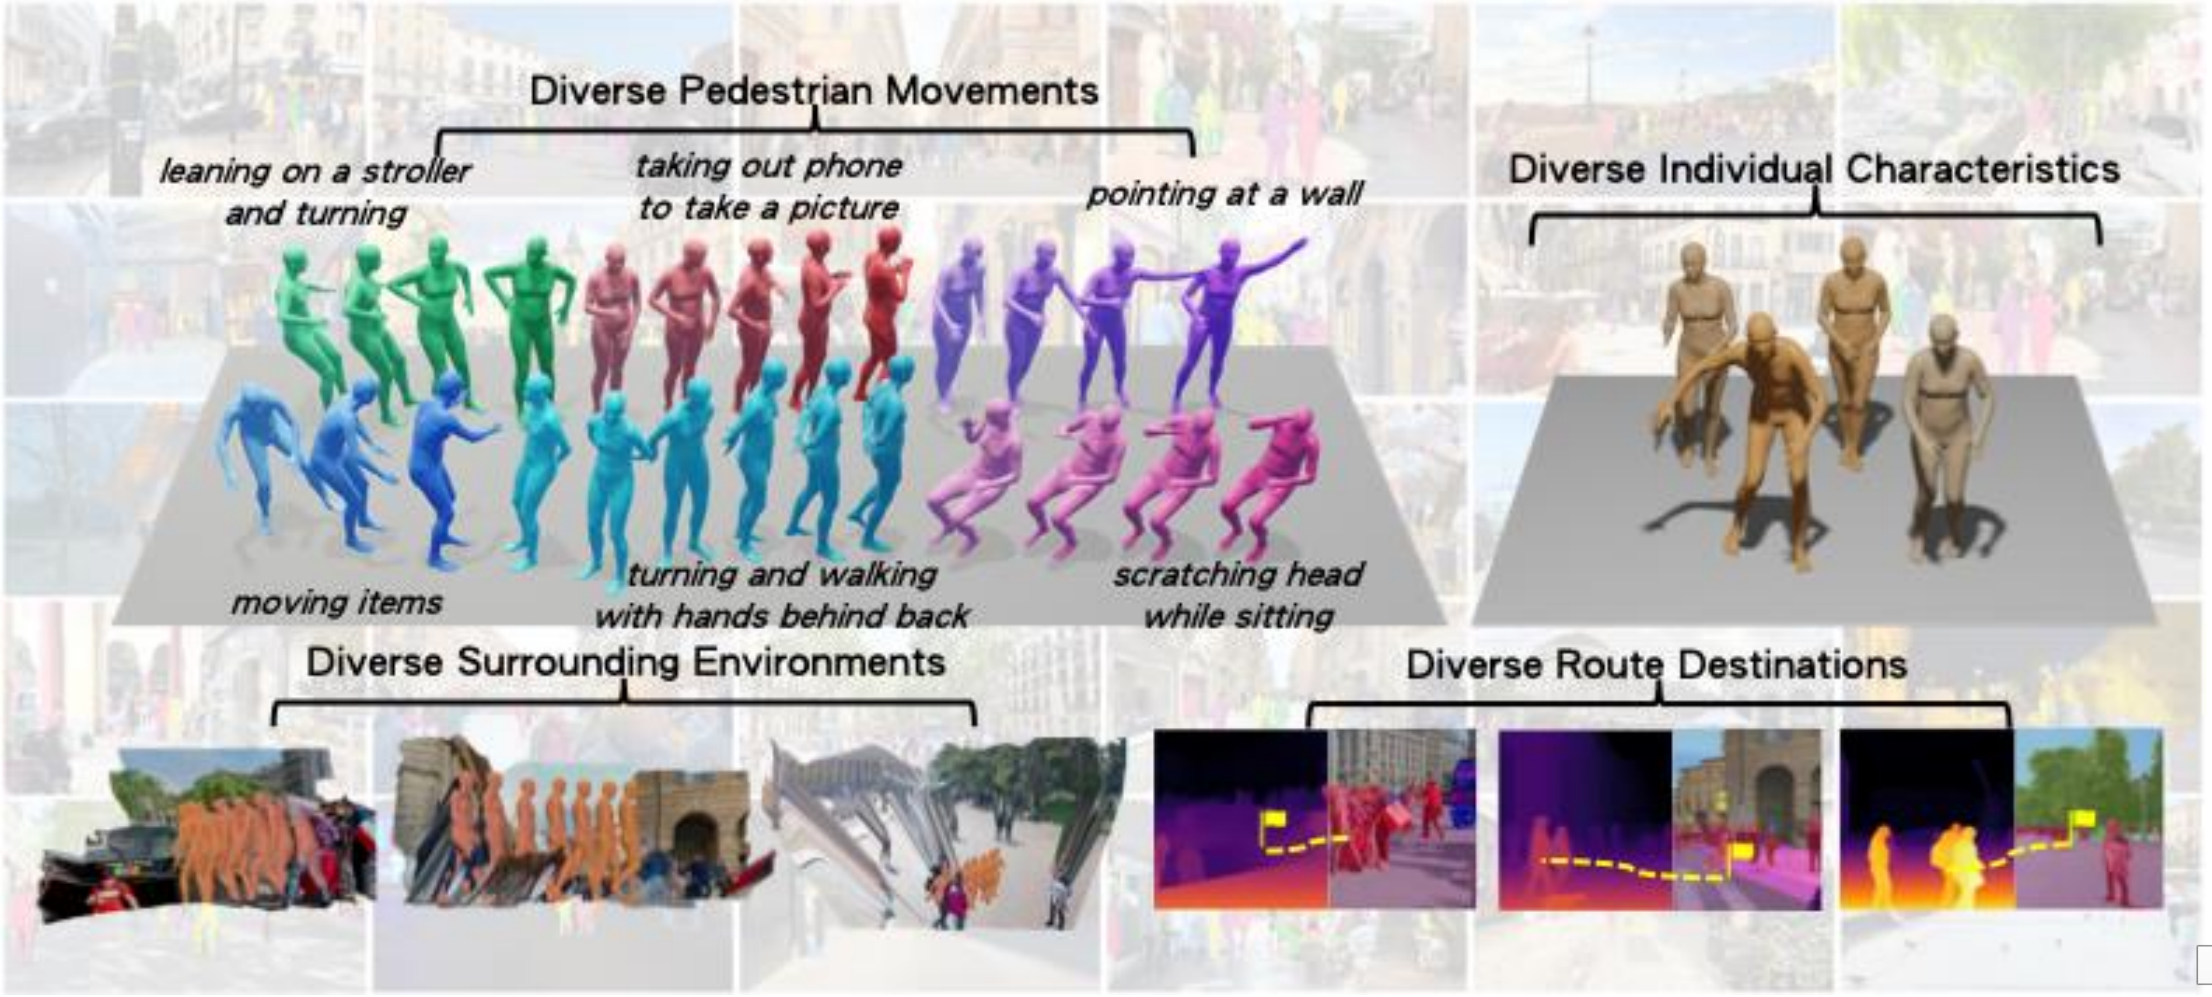

PedGen: Generative Model for Context-Aware Pedestrian Motion

- Diffusion model conditioned on scene context
  - Note: Same as image diffusion models from lecture. We add noise and denoise a **motion vector** instead of an image.

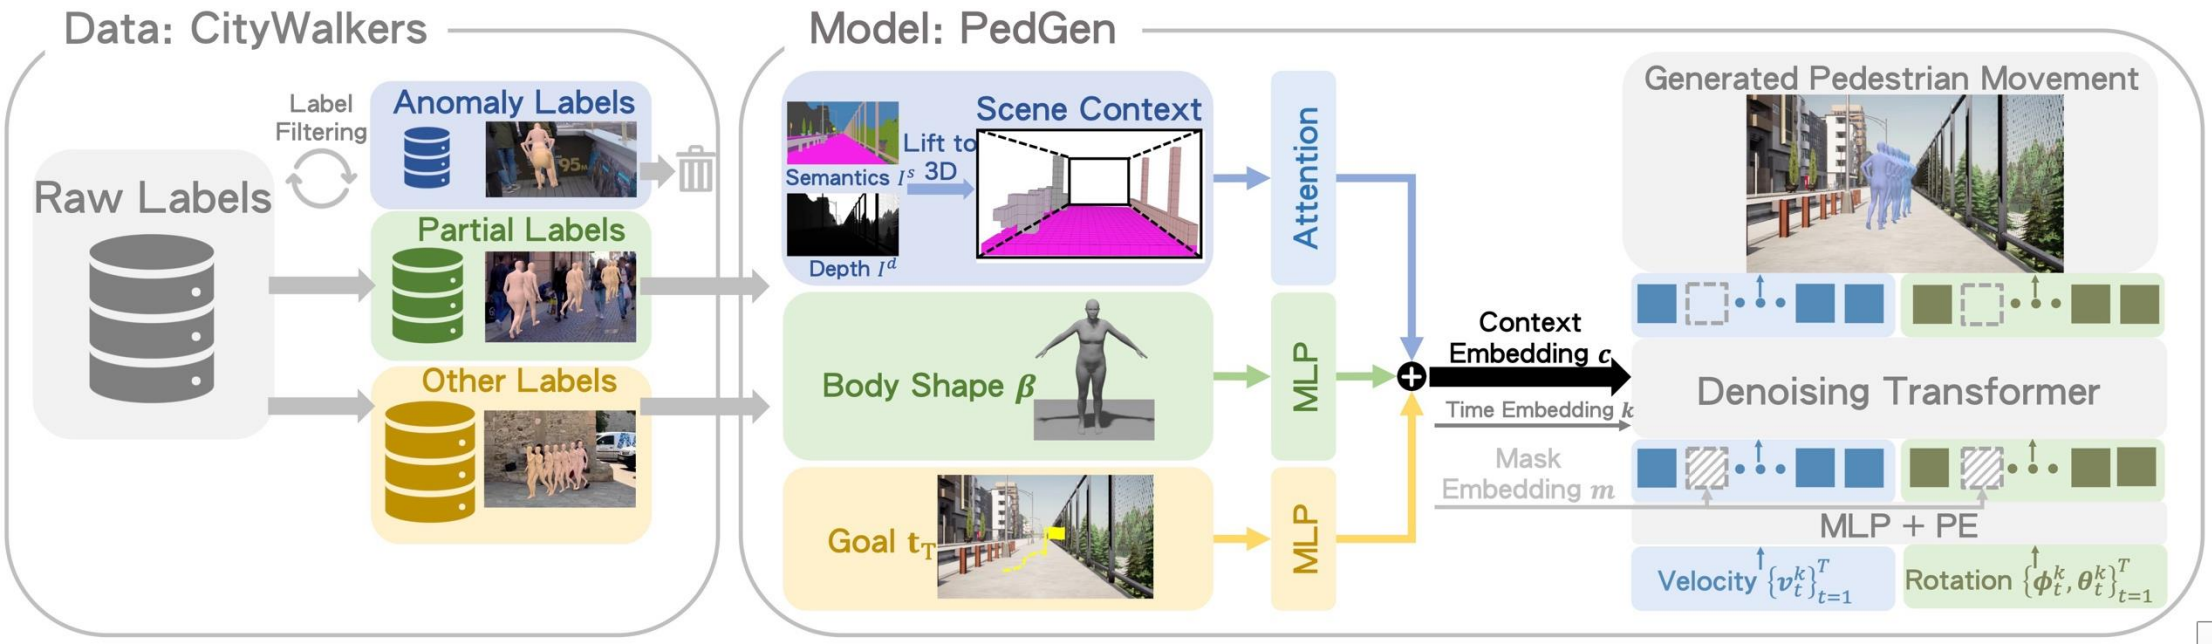

Generation Results in Real-World Scenarios and CARLA Simulation

Qualitative Comparison of Training PedGen with Scene Context Factors

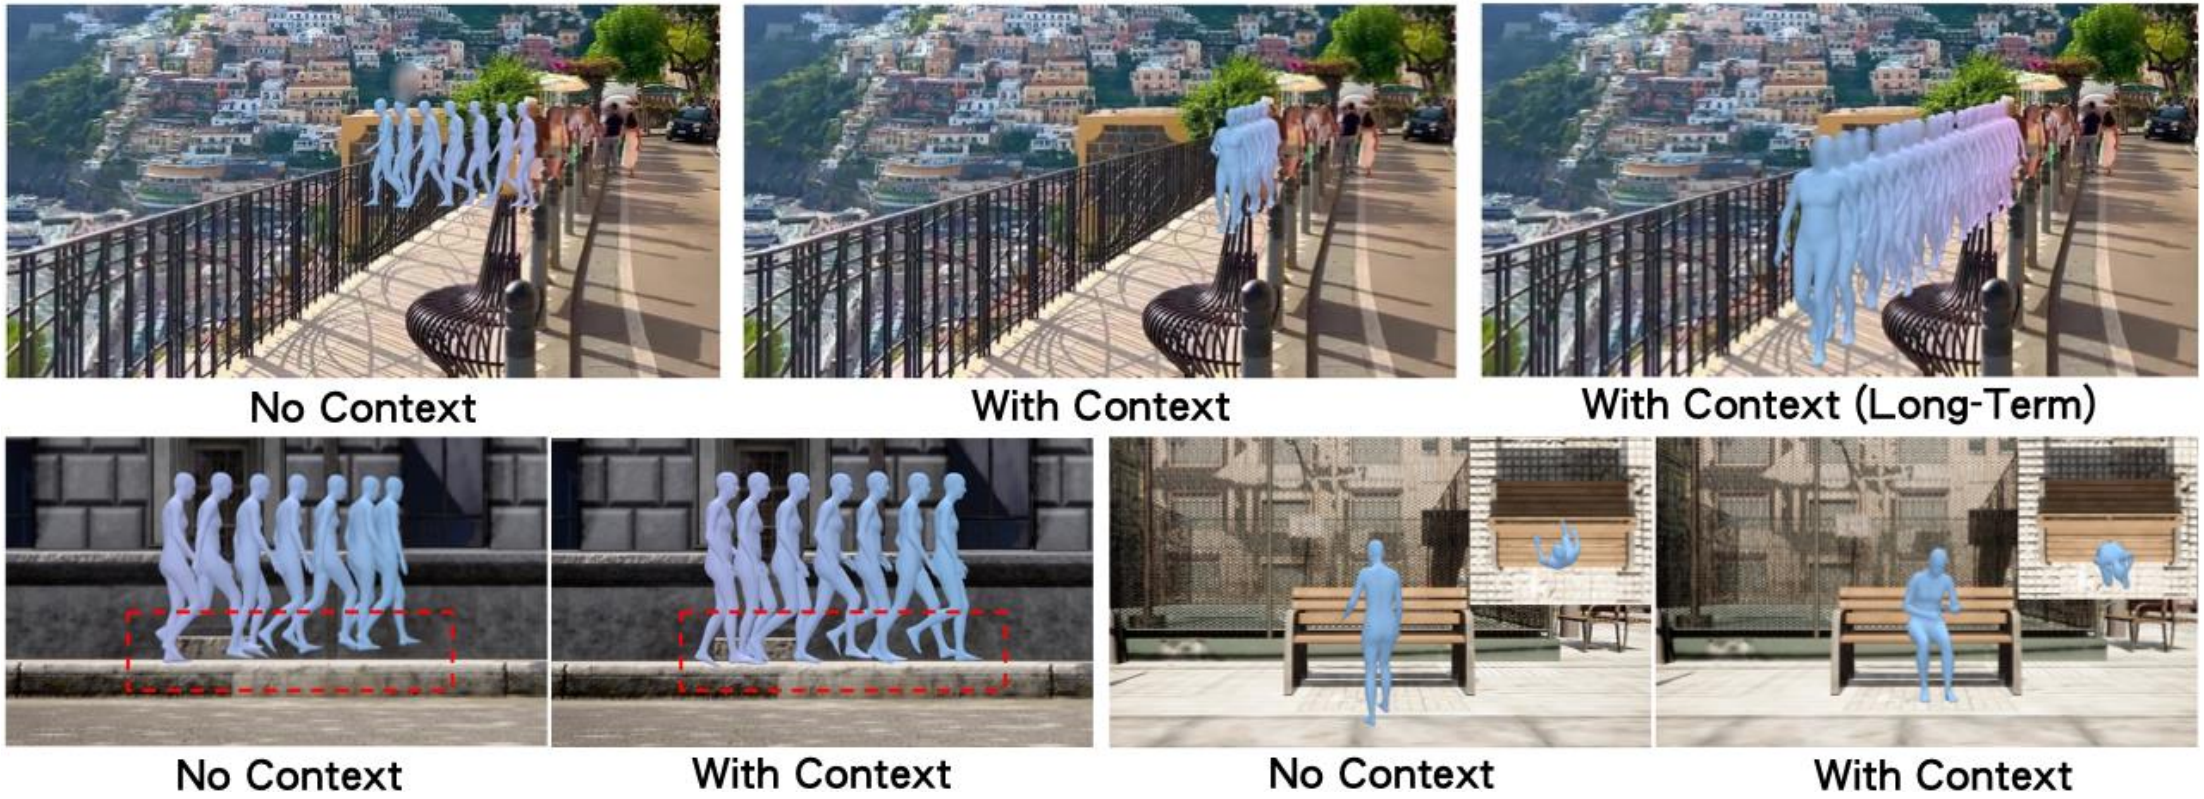

Q: How do we improve the pseudo-label quality?

A: JOSH

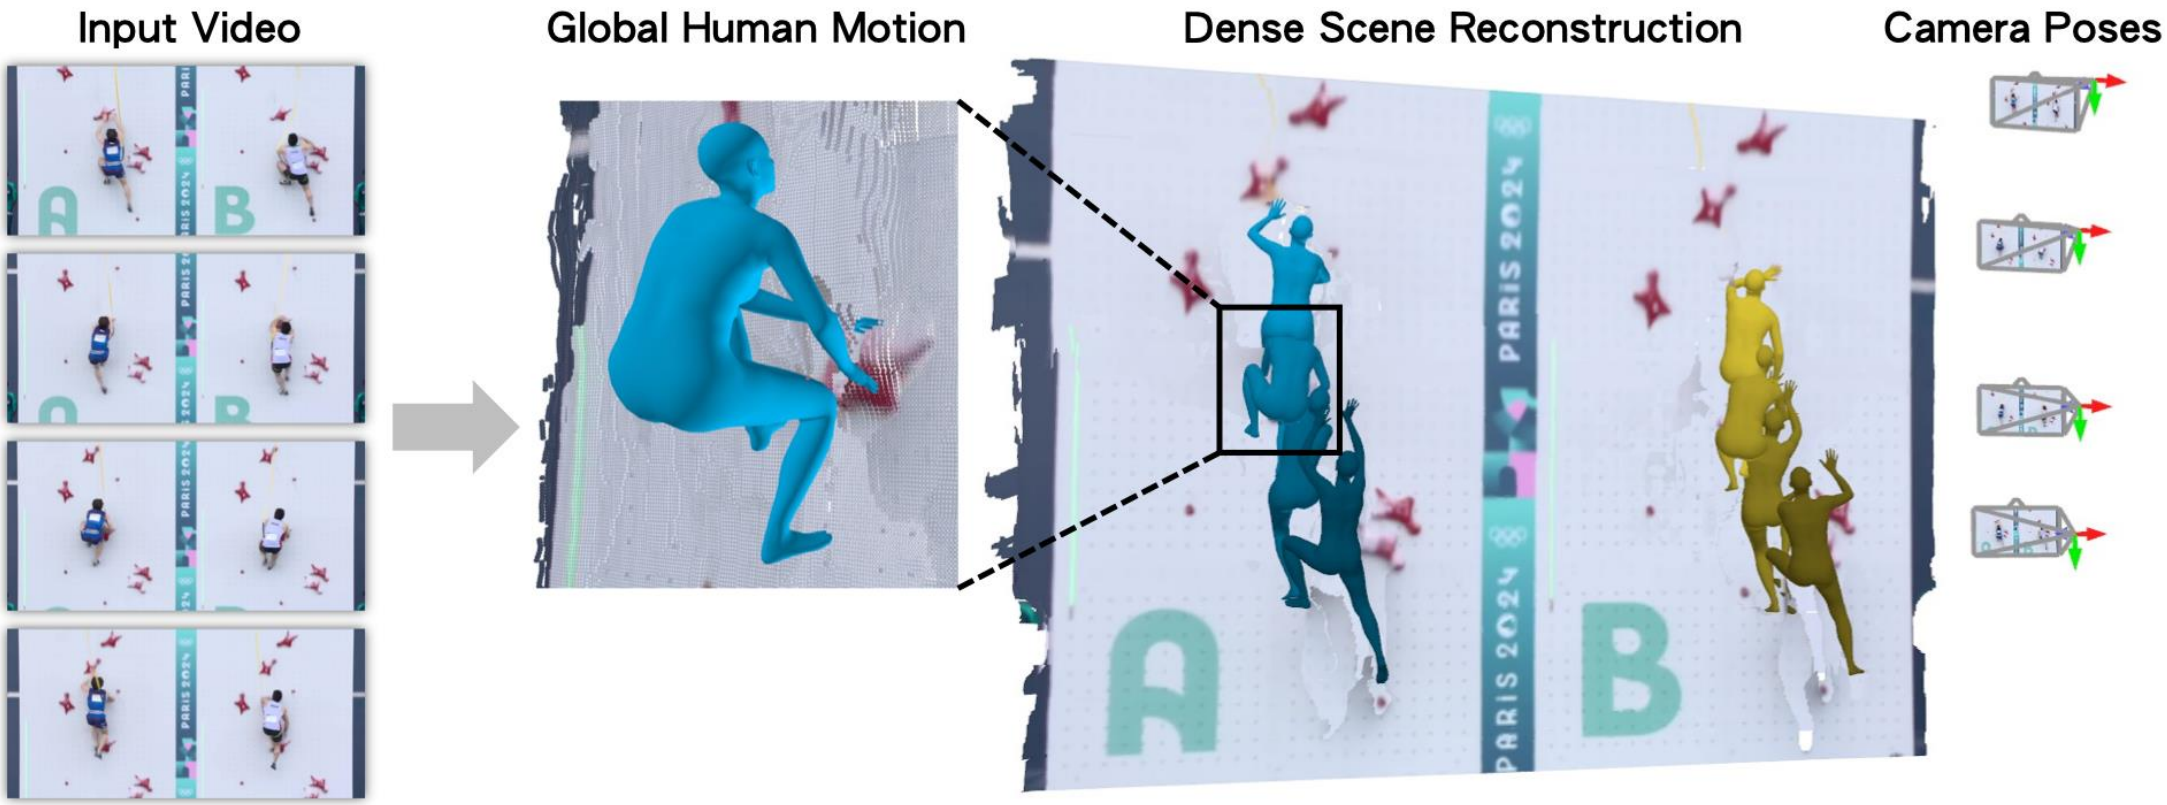

Joint Optimization for 4D Human-Scene Reconstruction in the Wild

Zhizheng Liu, Joe Lin, Wayne Wu, Bolei Zhou

Human Motion Estimation

- Predicting human movements (in SMPL) given a sequence of images
  - **TRAM** – Combines SLAM and depth estimator to predict metric-scale global trajectory

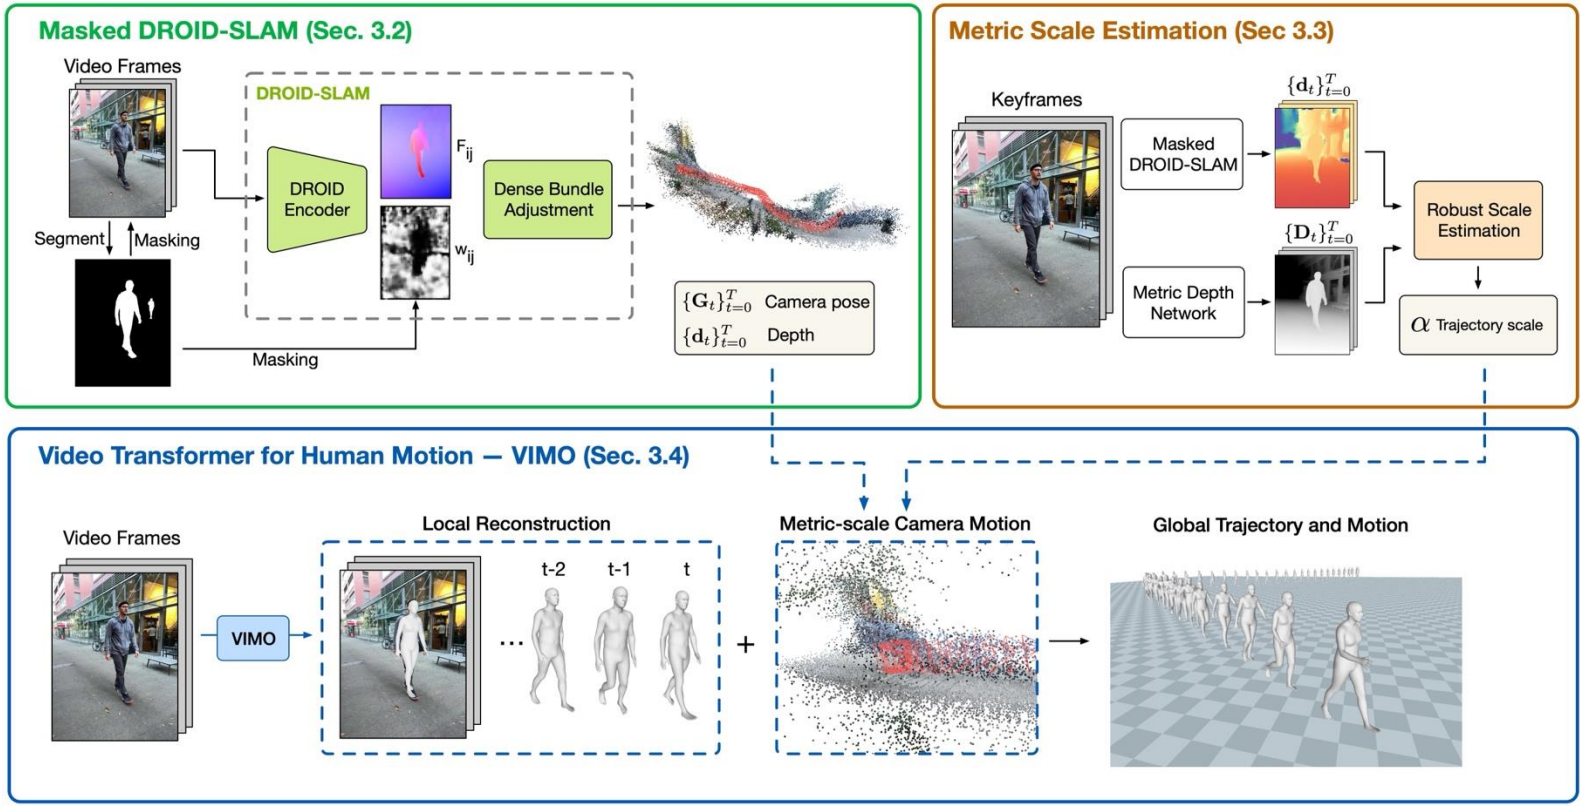

TRAM: Global Trajectory and Motion of 3D Humans from in-the-wild Videos. Arxiv 2024.

Scene Reconstruction

- Reconstructing the 3D scene given a set of images
  - Note: Most of the time we don't have access to camera intrinsics and extrinsics
- Classical CV approach
  - **Structure from Motion** (SfM) – Estimate camera poses and triangulate points with **stereo** and **epipolar** geometry
  - **COLMAP** – Dense reconstruction using SfM and bundle adjustment
  - **Problem: Computationally expensive and inaccurate with sparse scene views**

Scene Reconstruction

- Deep learning approach
  - **DUSt3R** – Dense and unconstrained stereo (two cameras) 3D reconstruction
  - **MASt3R** – Learning additional feature descriptor for efficient 3D point matching
  - **MASt3R-SfM** – Extending pipeline to videos and optimize dense depth, camera intrinsics, camera pose, global scale

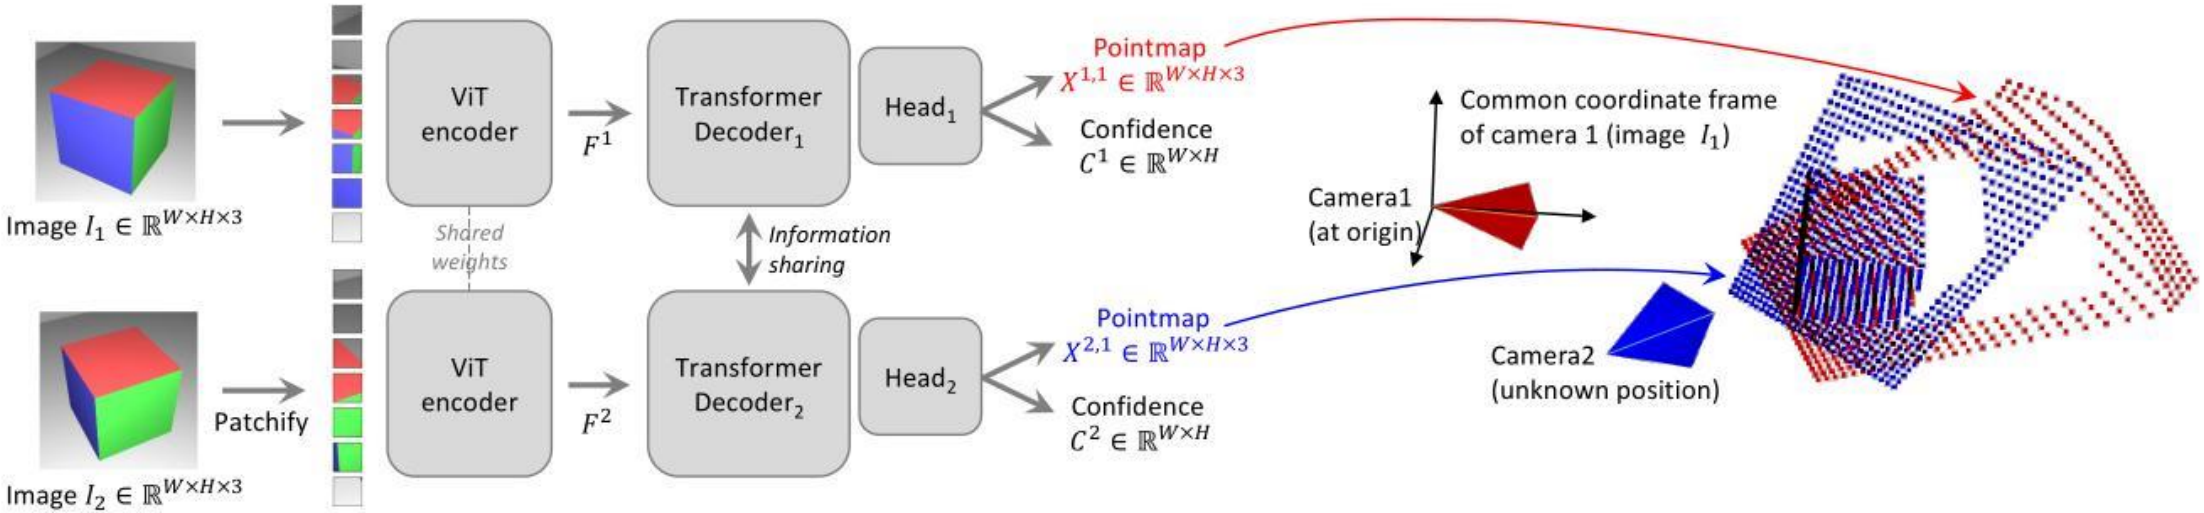

DUSt3R: Geometric 3D Vision Made Easy. CVPR 2024.

Scene Reconstruction

- Deep learning approach
  - **DUSt3R** – Dense and unconstrained stereo (two cameras) 3D reconstruction
  - **MASt3R** – Learning additional feature descriptor for efficient 3D point matching
  - **MASt3R-SfM** – Extending pipeline to videos and optimize dense depth, camera intrinsics, camera pose, global scale

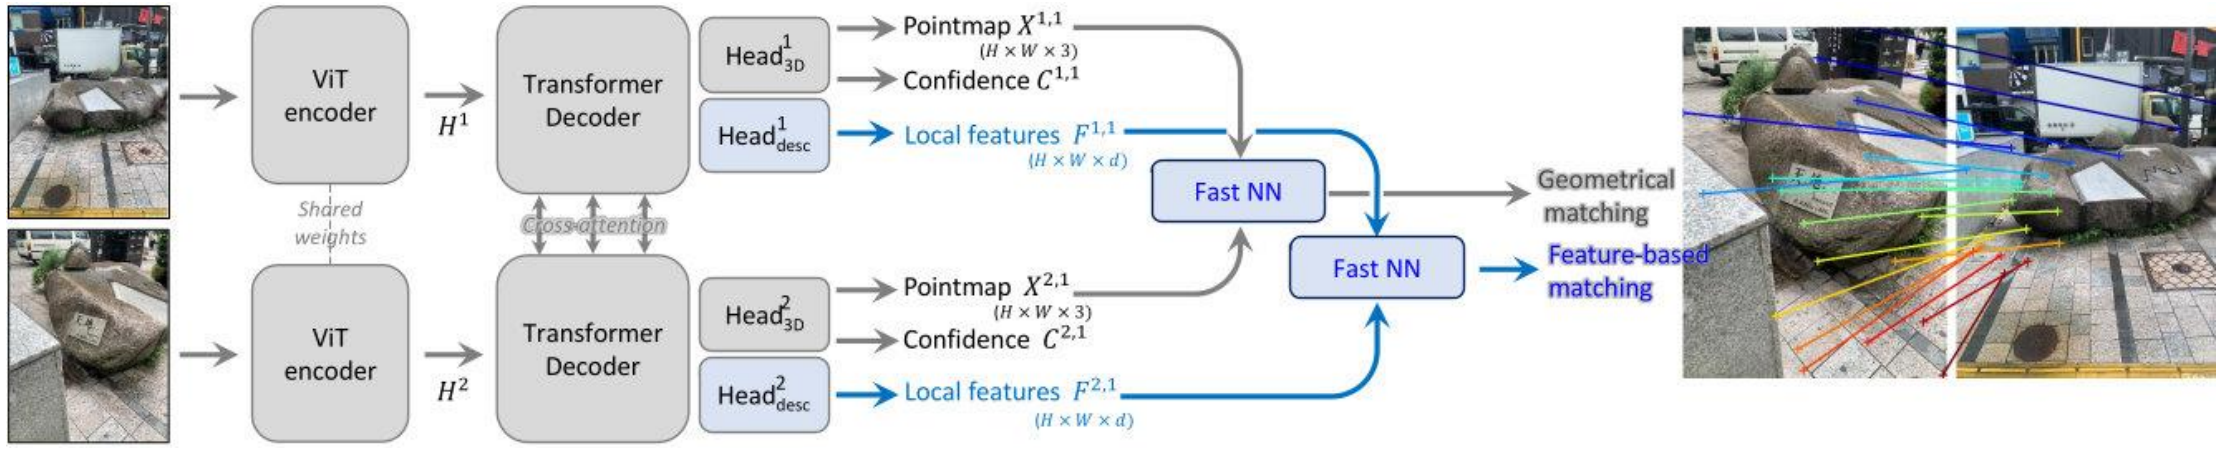

Grounding Image Matching in 3D with MASt3R. Arxiv 2024.

JOSH: Joint Optimization of Scene Geometry and Human Motion

- Initialize with MASt3R and TRAM + obtain pseudo-depths
- Add human into joint optimization
  - Contact Scale Loss $\mathcal{L}_{c1}$ – matches human-scene contact points
  - Contact Static Loss $\mathcal{L}_{c2}$ – ensures contact points remain stationary across time

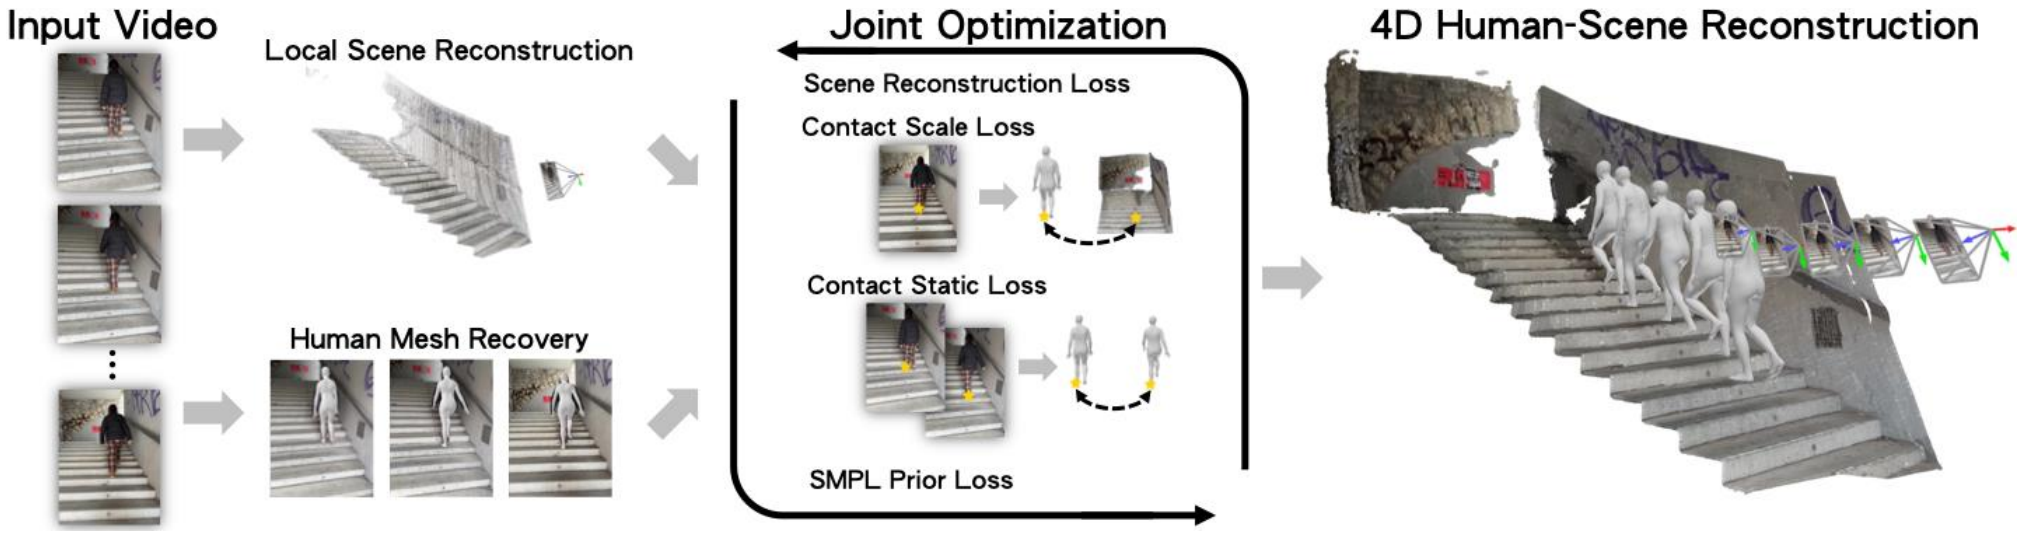

Evaluation Results for Global Human Motion on EMDB Dataset

Input Video from EMDB

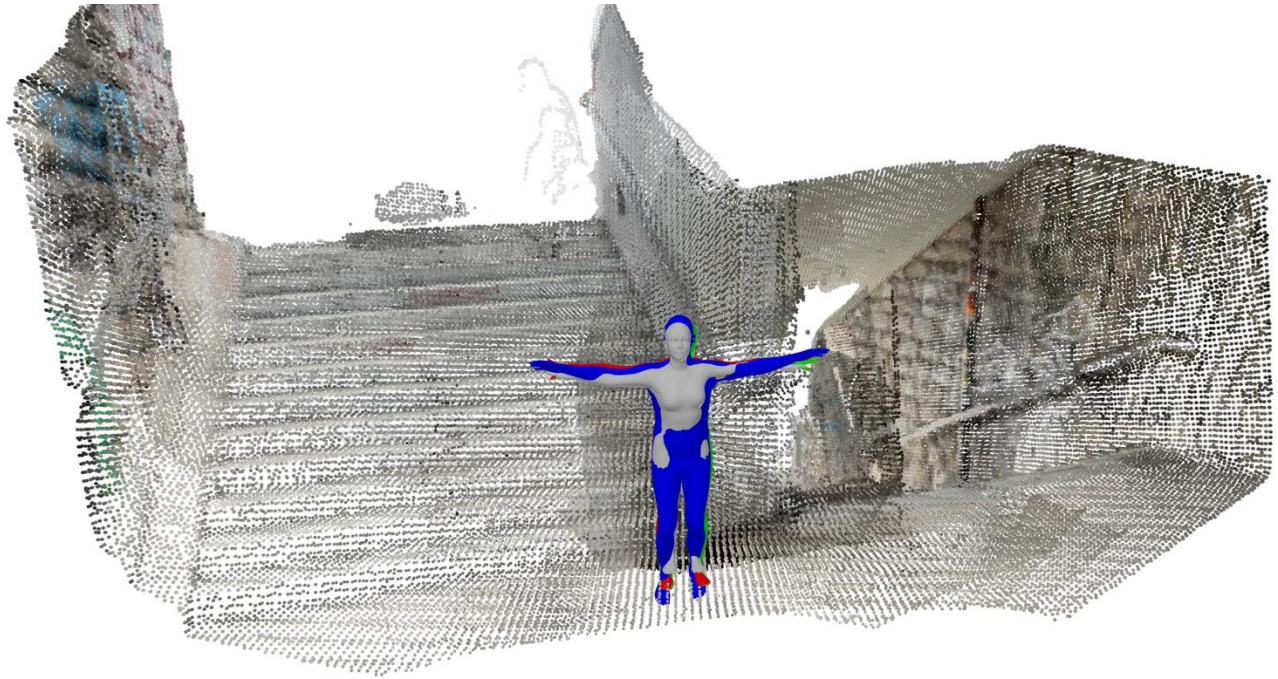

Ground Truth JOSH (ours) TRAM WHAM

*Dense scene reconstruction result is from JOSH as a reference, as no GT is available for EMDB and WHAM and TRAM cannot reconstruct the scene

Evaluation Results for Global Human Motion on SLOPER4D Dataset

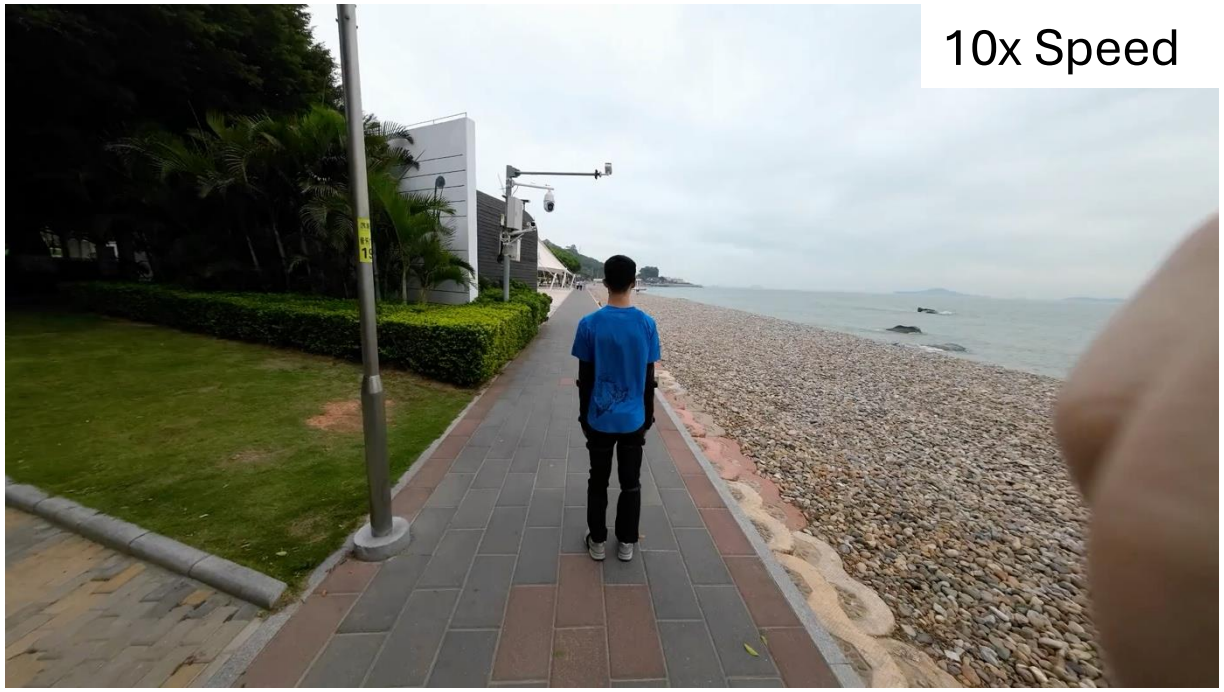

10x Speed

Input Video from SLOPER4D

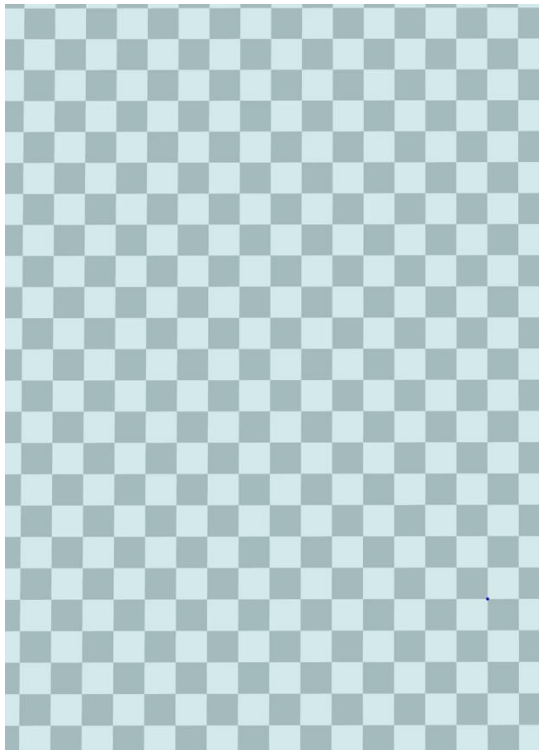

Ground Truth

JOSH (ours)

MASt3R

MonST3R

*This is a long sequence, and the checkerboard tile size is 10m x 10m

Human and Scene Reconstruction Results on in-the-wild Videos

Input Web Video

4D Human-Scene Reconstruction

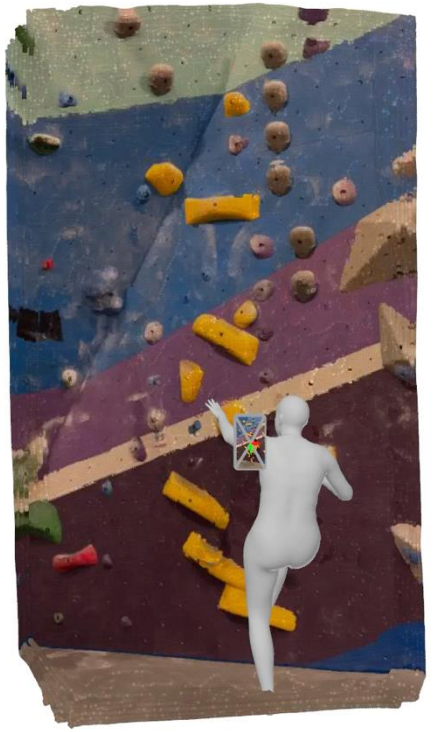

Thank you for listening!

If any of this interests you, feel free to reach out to Zhizheng Liu.